

Kernels

Course of Machine Learning
Master Degree in Computer Science
University of Rome "Tor Vergata"

Giorgio Gambosi

a.a. 2018-2019

- Thus far, we have been assuming that each object that we deal with can be represented as a fixed-size feature vector $\mathbf{x} \in \mathbb{R}^d$
- For certain kinds of objects (text document, protein sequence, parse tree, etc.) it is not clear how to best represent them in this way
 1. first approach: define a generative model of data (with latent variables) and define an object as the inferred values of latent variables
 2. second approach: do not rely on vector representation, but just assume a similarity measure between objects is defined


Representation by pairwise comparison

Idea

- Define a comparison function $\kappa : \mathcal{X} \times \mathcal{X} \mapsto \mathbb{R}$
- Represent a set of data items $\mathbf{x}_1, \dots, \mathbf{x}_n$ by the $n \times n$ Gram matrix \mathbf{G} such that

$$\mathbf{G}_{ij} = \kappa(\mathbf{x}_i, \mathbf{x}_j)$$

- \mathbf{G} is always an $n \times n$ matrix, whatever the nature of data: the same algorithm will work for any type of data (vectors, strings, ...)



Given a set χ , a function $\kappa : \chi^2 \mapsto \mathbb{R}$ is a **kernel** on χ if there exists a Hilbert space \mathcal{H} (essentially, a vector space with dot product \cdot) and a map $\phi : \chi \mapsto \mathcal{H}$ such that for all $\mathbf{x}_1, \mathbf{x}_2 \in \chi$ we have

$$\kappa(\mathbf{x}_1, \mathbf{x}_2) = \phi(\mathbf{x}_1) \cdot \phi(\mathbf{x}_2)$$

We shall consider the particular but common case when $\mathcal{H} = \mathbb{R}^d$ for some $d > 0$, $\phi(\mathbf{x}) = (\phi_1(\mathbf{x}), \dots, \phi_d(\mathbf{x}))$ and

$$\phi(\mathbf{x}_1) \cdot \phi(\mathbf{x}_2) = \phi(\mathbf{x}_1)^T \phi(\mathbf{x}_2)$$

ϕ is called a **feature map** a \mathcal{H} a **feature space** of κ

Positive definiteness of κ is a relevant property in this framework.

Positive semidefiniteness

Given a set χ , a function $\kappa : \chi^2 \mapsto \mathbb{R}$ is positive semidefinite if for all $n \in \mathbb{N}$, $(\mathbf{x}_1, \dots, \mathbf{x}_n) \in \chi^n$ the corresponding Gram matrix is positive semidefinite, that is $\mathbf{z}^T \mathbf{G} \mathbf{z} \geq 0$ for all vectors $\mathbf{z} \in \mathbb{R}^n$

Why is positive semidefiniteness relevant?

Let $\kappa : \mathcal{X} \times \mathcal{X} \mapsto \mathbb{R}$. Then κ is a kernel iff for all sets $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\}$ the corresponding Gram matrix \mathbf{G} is symmetric and positive semidefinite

Only if: $\mathbf{G}_{ij} = \phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j)$ then clearly $\mathbf{G}_{ij} = \mathbf{G}_{ji}$. Moreover for any $\mathbf{z} \in \mathbb{R}^d$

$$\begin{aligned}\mathbf{z}^T \mathbf{G} \mathbf{z} &= \sum_{i=1}^d \sum_{j=1}^d z_i \mathbf{G}_{ij} z_j = \sum_{i=1}^d \sum_{j=1}^d z_i \phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j) z_j \\&= \sum_{i=1}^d \sum_{j=1}^d z_i \left(\sum_{k=1}^d \phi_k(\mathbf{x}_i) \phi_k(\mathbf{x}_j) \right) z_j = \sum_{k=1}^n \sum_{i=1}^d \sum_{j=1}^d z_i \phi_k(\mathbf{x}_i) \phi_k(\mathbf{x}_j) z_j \\&= \sum_{k=1}^d \left(\sum_{i=1}^d z_i \phi_k(\mathbf{x}_i) \right)^2 \geq 0\end{aligned}$$

Why are positive definite kernels relevant?

If: Given $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\}$ if \mathbf{G} is positive definite it is possible to compute an eigenvector decomposition

$$\mathbf{G} = \mathbf{U}^T \mathbf{\Lambda} \mathbf{U}$$

where $\mathbf{\Lambda}$ is the diagonal matrix of eigenvalues $\lambda_i > 0$ and the columns $\mathbf{u}_1, \dots, \mathbf{u}_n$ of \mathbf{U} are the corresponding eigenvectors. Then,

$$\mathbf{G}_{ij} = (\mathbf{\Lambda}^{\frac{1}{2}} \mathbf{u}_i)^T (\mathbf{\Lambda}^{\frac{1}{2}} \mathbf{u}_j)$$

Then if we define $\phi(\mathbf{x}_i) = \mathbf{\Lambda}^{\frac{1}{2}} \mathbf{u}_i$ we get

$$\kappa(\mathbf{x}_i, \mathbf{x}_j) = \phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j) = \mathbf{G}_{ij}$$

This results is valid only wrt the domain $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\}$. For the general case, consider $n \rightarrow \infty$ (as for example in gaussian processes)

Why are positive definite kernels relevant?

Using positive definite kernels allows to apply the **kernel trick** wherever useful.

Kernel trick

Any algorithm which processes finite-dimensional vectors in such a way to consider only pairwise dot products can be applied to higher (possibly infinite) dimensional vectors by replacing each dot product by a suitable application of a positive definite kernel.

- Many practical applications
- Vectors in the new space are manipulated only implicitly, through pairwise dot products, computed by evaluating the kernel function on the original pair of vectors

Example: Support vector machines. Also, many linear models for regression and classification can be reformulated in terms of a **dual representation**

Dual representations: example

Regularized sum of squares in regression with predefined basis function $\phi(\mathbf{x})$

$$\begin{aligned} J(\mathbf{w}) &= \frac{1}{2} \sum_{i=1}^n \left(\mathbf{w}^T \phi(\mathbf{x}_i) - t_i \right)^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w} \\ &= \frac{1}{2} (\Phi \mathbf{w} - \mathbf{t})^T (\Phi \mathbf{w} - \mathbf{t}) + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w} \end{aligned}$$

where by definition of $\Phi \in \mathbb{R}^{n \times d}$ it is $\Phi_{ij} = \phi_j(\mathbf{x}_i)$

Setting $\frac{\partial J(\mathbf{w})}{\partial \mathbf{w}} = \mathbf{0}$, the resulting solution is

$$\hat{\mathbf{w}} = (\Phi \Phi^T + \lambda \mathbf{I}_d)^{-1} \Phi^T \mathbf{t} = \Phi^T (\Phi \Phi^T + \lambda \mathbf{I}_n)^{-1} \mathbf{t}$$

since it is possible to prove that for any matrix $\mathbf{A} \in \mathbb{R}^{r \times c}$ it is

$$(\mathbf{A}^T \mathbf{A} + \lambda \mathbf{I}_r)^{-1} \mathbf{A}^T = \mathbf{A}^T (\mathbf{A} \mathbf{A}^T + \lambda \mathbf{I}_c)^{-1}$$

Dual representations: example

If we define the **dual variables** $\mathbf{a} = (\Phi\Phi^T + \lambda\mathbf{I}_n)^{-1}\mathbf{t}$, we get $\mathbf{w} = \Phi^T\mathbf{a}$.

By substituting $\Phi^T\mathbf{a}$ to \mathbf{w} we express the cost function in terms of \mathbf{a} , instead of \mathbf{w} , introducing a **dual formulation** of J .

$$\begin{aligned} J(\mathbf{a}) &= \frac{1}{2}\mathbf{a}^T\Phi\Phi^T\Phi\Phi^T\mathbf{a} + \frac{1}{2}\mathbf{t}^T\mathbf{t} - \mathbf{a}^T\Phi\Phi^T\mathbf{t} + \frac{\lambda}{2}\mathbf{a}^T\Phi\Phi^T\mathbf{a} \\ &= \frac{1}{2}\mathbf{a}^T\mathbf{G}\mathbf{G}\mathbf{a} + \frac{1}{2}\mathbf{t}^T\mathbf{t} - \mathbf{a}^T\mathbf{G}\mathbf{t} + \frac{\lambda}{2}\mathbf{a}^T\mathbf{G}\mathbf{a} \end{aligned}$$

where $\mathbf{G} = \Phi\Phi^T$ is the Gram matrix, such that by definition

$$\mathbf{G}_{ij} = \sum_{k=1}^d \phi_k(\mathbf{x}_i)\phi_k(\mathbf{x}_j) = \phi(\mathbf{x}_i)^T\phi(\mathbf{x}_j)$$

Dual representations: example

Setting the gradient of $\frac{\partial J(\mathbf{a})}{\partial \mathbf{a}} = \mathbf{0}$ it results

$$\hat{\mathbf{a}} = (\mathbf{G} + \mathbf{I}\lambda_n)^{-1}\mathbf{t}$$

We can use this to make predictions in a different way

$$\begin{aligned}y(\mathbf{x}) &= \mathbf{w}^T \phi(\mathbf{x}) = \mathbf{a}^T \Phi \phi(\mathbf{x}) = \mathbf{t}^T (\mathbf{G} + \mathbf{I}\lambda_n)^{-1} \Phi \phi(\mathbf{x}) \\&= \mathbf{k}(\mathbf{x}) (\mathbf{G} + \mathbf{I}\lambda_n)^{-1} \mathbf{t}\end{aligned}$$

where

$$\begin{aligned}\mathbf{k}(\mathbf{x}) &= \phi(\mathbf{x})^T \Phi = (\phi(\mathbf{x}_1)^T \phi(\mathbf{x}), \dots, \phi(\mathbf{x}_n^T \phi(\mathbf{x}))^T \\&= (\kappa(\mathbf{x}_1, \mathbf{x}), \dots, \kappa(\mathbf{x}_n, \mathbf{x}))^T \\&= (\kappa_1(\mathbf{x}), \dots, \kappa_n(\mathbf{x}))^T\end{aligned}$$

Dual representations: another example

- As well known, a perceptron is a linear classifier with prediction $y(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$
- Its update rule is: If \mathbf{x}_i is misclassified, that is $\mathbf{w}^T \mathbf{x}_i t_i < 0$, then $\mathbf{w} := \mathbf{w} + t_i \mathbf{x}_i$
- If we assume a zero initial value for all w_k , then \mathbf{w} is the sum of all items that have been considered as misclassified by the algorithm, where each item is weighted by the number of times it has been considered.
- We may then define a **dual** formulation by setting $\mathbf{w} = \sum_{k=1}^n a_k \mathbf{x}_k$, which results in prediction $y(\mathbf{x}) = \sum_{k=1}^n a_k \mathbf{x}_k^T \mathbf{x}$
- and update rule: if \mathbf{x}_i is misclassified, that is $\sum_{k=1}^n a_k \mathbf{x}_k^T \mathbf{x}_i < 0$, then $a_i := a_i + 1$
- a **kernelized** perceptron can be defined with $y(\mathbf{x}) = \sum_{k=1}^n a_k \phi(\mathbf{x}_k)^T \phi(\mathbf{x})$ or with $y(\mathbf{x}) = \sum_{k=1}^n a_k \kappa(\mathbf{x}_k, \mathbf{x})$, by just using a positive definite kernel κ

Kernelization: one more example

- The k -nn classifier selects the label of the nearest neighbor: assume the Euclidean distance is considered

$$\|\mathbf{x}_i - \mathbf{x}_j\|^2 = \mathbf{x}_i^T \mathbf{x}_i + \mathbf{x}_j^T \mathbf{x}_j - 2\mathbf{x}_i^T \mathbf{x}_j$$

- We can now replace the dot products by a valid positive definite kernel and we obtain:

$$d(\mathbf{x}_i, \mathbf{x}_j)^2 = \kappa(\mathbf{x}_i, \mathbf{x}_i) + \kappa(\mathbf{x}_j, \mathbf{x}_j) - 2\kappa(\mathbf{x}_i, \mathbf{x}_j)$$

- This is a kernelized nearest-neighbor classifier
- We do not explicitly compute vectors

Why referring to the dual representation?

- While in the original formulation of linear regression \mathbf{w} can be derived by inverting the $m \times m$ matrix $\Phi^T \Phi$, in the dual formulation computing \mathbf{a} requires inverting the $n \times n$ matrix $\mathbf{G} + \mathbf{I}\lambda$.
- Since usually $n \gg m$, this seems to lead to a loss of efficiency.

Since not all functions $f : \mathcal{X} \mapsto \mathbb{R}^d$ are positive definite kernel, some method to define them must be applied.

- the straightforward way is just to define a basis function ϕ and define $\kappa(\mathbf{x}_1, \mathbf{x}_2) = \phi(\mathbf{x}_1)^T \phi(\mathbf{x}_2)$. κ is a positive definite kernel since

1. $\phi(\mathbf{x}_1)^T \phi(\mathbf{x}_2) = \phi(\mathbf{x}_2)^T \phi(\mathbf{x}_1)$
2. $\sum_{i=1}^n \sum_{j=1}^n c_i c_j \kappa(\mathbf{x}_i, \mathbf{x}_j) = \sum_{i=1}^n \sum_{j=1}^n c_i c_j \phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j) = \left\| \sum_{i=1}^n c_i \phi(\mathbf{x}_i) \right\|^2 \geq 0$

- a second method defines a possible kernel function κ directly: in order to ensure that such function is a valid kernel, apply Mercer's theorem and prove that κ is a positive definite kernel by showing it is symmetric and the corresponding Gram matrix \mathbf{G} is positive definite for all possible sets of items. In this case we do not define ϕ

A simple positive definite kernel

Let $\chi = \mathbb{R}$: the function $\kappa : \mathbb{R}^2 \mapsto \mathbb{R}$ defined as

$$\kappa(x_1, x_2) = x_1 x_2$$

is a positive definite kernel. In fact,

- $x_1 x_2 = x_2 x_1$
- $\sum_{i=1}^n \sum_{j=1}^n c_i c_j \kappa(x_i, x_j) = \sum_{i=1}^n \sum_{j=1}^n c_i c_j x_i x_j = \left(\sum_{i=1}^n c_i x_i \right)^2 \geq 0$

Another simple positive definite kernel

Let $\chi = \mathbb{R}^d$: the function $\kappa : \chi^2 \mapsto \mathbb{R}$ defined as

$$\kappa(\mathbf{x}_1, \mathbf{x}_2) = \mathbf{x}_1^T \mathbf{x}_2$$

is a positive definite kernel. In fact,

$$\begin{aligned} & \cdot \mathbf{x}_1^T \mathbf{x}_2 = \mathbf{x}_2^T \mathbf{x}_1 \\ & \cdot \sum_{i=1}^n \sum_{j=1}^n c_i c_j \kappa(\mathbf{x}_i, \mathbf{x}_j) = \sum_{i=1}^n \sum_{j=1}^n c_i c_j \mathbf{x}_i^T \mathbf{x}_j = \\ & \quad \left\| \sum_{i=1}^n c_i \mathbf{x}_i \right\|^2 \geq 0 \end{aligned}$$

- a third method defines again a possible kernel function κ directly: in order to ensure that such function is a valid kernel, a basis function ϕ must be found such that $\kappa(\mathbf{x}_1, \mathbf{x}_2) = \phi(\mathbf{x}_1)^T \phi(\mathbf{x}_2)$ for all $\mathbf{x}_1, \mathbf{x}_2$

Example

A polynomial kernel in 2d: $\phi(\mathbf{x}) = (x_1^2, \sqrt{2}x_1x_2, x_2^2)$

$$\begin{aligned}\kappa(\mathbf{x}_1, \mathbf{x}_2) &= (x_{11}^2, \sqrt{2}x_{11}x_{12}, x_{12}^2)^T (x_{21}^2, \sqrt{2}x_{21}x_{22}, x_{22}^2) \\ &= x_{11}^2x_{21}^2 + 2x_{11}x_{12}x_{21}x_{22} + x_{12}^2x_{22}^2 \\ &= \|\mathbf{x}_1^T \mathbf{x}_2\|^2\end{aligned}$$

Example

If $\mathbf{x}_1, \mathbf{x}_2 \in \mathbb{R}^d$ define $\kappa(\mathbf{x}_1, \mathbf{x}_2) = (\mathbf{x}_1 \cdot \mathbf{x}_2)^2 = \phi(\mathbf{x}_1)^T \phi(\mathbf{x}_2)$, where

$$\phi(\mathbf{x}) = (x_1^2, \dots, x_d^2, x_1x_2, \dots, x_1x_d, x_2x_1, \dots, x_dx_{d-1})^T$$

Example

$\kappa(\mathbf{x}_1, \mathbf{x}_2) = (\mathbf{x}_1 \cdot \mathbf{x}_2)^2$ is a valid kernel function, since

$$\begin{aligned}\kappa(\mathbf{x}_1, \mathbf{x}_2) &= (x_{11}x_{21} + x_{12}x_{22})^2 \\&= x_{11}^2x_{21}^2 + x_{12}^2x_{22}^2 + 2x_{11}x_{12}x_{21}x_{22} \\&= (x_{11}^2, x_{12}^2, x_{11}x_{12}, x_{11}x_{12}) \cdot (x_{21}^2, x_{22}^2, x_{21}x_{22}, x_{21}x_{22}) \\&= \phi(\mathbf{x}_1) \cdot \phi(\mathbf{x}_2)\end{aligned}$$

The basis function thus results $\phi(\mathbf{x}) = (x_1^2, x_2^2, x_1x_2, x_1x_2)^T$.

Example

- In general, if $\mathbf{x}_1, \mathbf{x}_2 \in \mathbb{R}^d$ then
 $\kappa(\mathbf{x}_1, \mathbf{x}_2) = (\mathbf{x}_1 \cdot \mathbf{x}_2)^2 = \boldsymbol{\phi}(\mathbf{x}_1)^T \boldsymbol{\phi}(\mathbf{x}_2)$, where

$$\boldsymbol{\phi}(\mathbf{x}) = (x_1^2, \dots, x_d^2, x_1x_2, \dots, x_1x_d, x_2x_1, \dots, x_dx_{d-1})^T$$

- the d -dimensional input space is mapped onto a space with dimension $m = d^2$
- observe that computing $\kappa(\mathbf{x}_1, \mathbf{x}_2)$ requires time $O(d)$, while deriving it from $\boldsymbol{\phi}(\mathbf{x}_1)^T \boldsymbol{\phi}(\mathbf{x}_2)$ requires $O(d^2)$ steps

Example

$\kappa(\mathbf{x}_1, \mathbf{x}_2) = (\mathbf{x}_1 \cdot \mathbf{x}_2 + c)^2$ is a kernel function, since

$$\begin{aligned}\kappa(\mathbf{x}_1, \mathbf{x}_2) &= (\mathbf{x}_1 \cdot \mathbf{x}_2 + c)^2 \\ &= \sum_{i=1}^n \sum_{j=1}^n x_{1i} x_{1j} x_{2i} x_{2j} + \sum_{i=1}^n (\sqrt{2c} x_{1i})(\sqrt{2c} x_{2i}) + c^2 \\ &= \phi(\mathbf{x}_1)^T \phi(\mathbf{x}_2)\end{aligned}$$

for

$$\phi(\mathbf{x}) = (x_1^2, \dots, x_d^2, x_1 x_2, \dots, x_1 x_d, x_2 x_1, \dots, x_d x_{d-1}, \sqrt{2c} x_1, \dots, \sqrt{2c} x_d)$$

This implies a mapping from a d -dimensional to a $(d+1)^2$ -dimensional space.

Example

$\kappa(\mathbf{x}_1, \mathbf{x}_2) = (\mathbf{x}_1 \cdot \mathbf{x}_2 + c)^t$ is a kernel function corresponding to a mapping from a d -dimensional space to a space of dimension

$$m = \sum_{i=0}^t d^i = \frac{d^{t+1} - 1}{d - 1}$$

corresponding to all products $x_{i_1} x_{i_2} \dots x_{i_l}$ with $0 \leq l \leq t$.

Observe that, even if the space has dimension $O(d^t)$, evaluating the kernel function requires just time $O(d)$.

Constructing kernels from kernels

More complex kernels can be derived from simpler ones by applying suitable transformation and composition rules. In fact, given kernel functions $\kappa_1(\mathbf{x}_1, \mathbf{x}_2)$, $\kappa_2(\mathbf{x}_1, \mathbf{x}_2)$, the function $\kappa(\mathbf{x}_1, \mathbf{x}_2)$ is a kernel in all the following cases

- $\kappa(\mathbf{x}_1, \mathbf{x}_2) = e^{\kappa_1(\mathbf{x}_1, \mathbf{x}_2)}$
- $\kappa(\mathbf{x}_1, \mathbf{x}_2) = \kappa_1(\mathbf{x}_1, \mathbf{x}_2) + \kappa_2(\mathbf{x}_1, \mathbf{x}_2)$
- $\kappa(\mathbf{x}_1, \mathbf{x}_2) = \kappa_1(\mathbf{x}_1, \mathbf{x}_2)\kappa_2(\mathbf{x}_1, \mathbf{x}_2)$
- $\kappa(\mathbf{x}_1, \mathbf{x}_2) = c\kappa_1(\mathbf{x}_1, \mathbf{x}_2)$, for any $c > 0$
- $\kappa(\mathbf{x}_1, \mathbf{x}_2) = \mathbf{x}_1^T \mathbf{A} \mathbf{x}_2$, with \mathbf{A} positive definite
- $\kappa(\mathbf{x}_1, \mathbf{x}_2) = f(\mathbf{x}_1)\kappa_1(\mathbf{x}_1, \mathbf{x}_2)g(\mathbf{x}_2)$, for any $f, g : \mathbb{R}^n \mapsto \mathbb{R}$
- $\kappa(\mathbf{x}_1, \mathbf{x}_2) = p(\kappa_1(\mathbf{x}_1, \mathbf{x}_2))$, for any polynomial $p : \mathbb{R}^q \mapsto \mathbb{R}$ with non-negative coefficients
- $\kappa(\mathbf{x}_1, \mathbf{x}_2) = \kappa_3(\phi(\mathbf{x}_1), \phi(\mathbf{x}_2))$, for any vector ϕ of m functions $\phi_i : \mathbb{R}^n \mapsto \mathbb{R}$ and for any kernel function $\kappa_3(\mathbf{x}_1, \mathbf{x}_2)$ in \mathbb{R}^m

Costructing kernel functions

$\kappa(\mathbf{x}_1, \mathbf{x}_2) = (\mathbf{x}_1 \cdot \mathbf{x}_2 + c)^d$ is a kernel function. In fact,

1. $\mathbf{x}_1 \cdot \mathbf{x}_2 = \mathbf{x}_1^T \mathbf{x}_2$ is a kernel function corresponding to the base functions $\phi = (\phi_1, \dots, \phi_n)$, with $\phi_i(\mathbf{x}) = \mathbf{x}$
2. c is a kernel function corresponding to the base functions $\phi = (\phi_1, \dots, \phi_n)$, with $\phi_i(\mathbf{x}) = \frac{\sqrt{c}}{n}$
3. $\mathbf{x}_1 \cdot \mathbf{x}_2 + c$ is a kernel function since it is the sum of two kernel functions
4. $(\mathbf{x}_1 \cdot \mathbf{x}_2 + c)^d$ is a kernel function since it is a polynomial with non negative coefficients (in particular $p(z) = z^d$) of a kernel function

Costructing kernel functions

$$\kappa(\mathbf{x}_1, \mathbf{x}_2) = e^{-\frac{\|\mathbf{x}_1 - \mathbf{x}_2\|^2}{2\sigma^2}}$$

is a kernel function. In fact,

1. since $\|\mathbf{x}_1 - \mathbf{x}_2\|^2 = \mathbf{x}_1^T \mathbf{x}_1 + \mathbf{x}_2^T \mathbf{x}_2 - 2\mathbf{x}_1^T \mathbf{x}_2$, it results

$$\kappa(\mathbf{x}_1, \mathbf{x}_2) = e^{-\frac{\mathbf{x}_1^T \mathbf{x}_1}{2\sigma^2}} e^{-\frac{\mathbf{x}_2^T \mathbf{x}_2}{2\sigma^2}} e^{\frac{\mathbf{x}_1^T \mathbf{x}_2}{\sigma^2}}$$

2. $\mathbf{x}_1^T \mathbf{x}_2$ is a kernel function (see above)

3. then, $\frac{\mathbf{x}_1^T \mathbf{x}_2}{\sigma^2}$ is a kernel function, being the product of a kernel

function with a constant $c = \frac{1}{\sigma^2}$

4. $e^{\frac{\mathbf{x}_1^T \mathbf{x}_2}{\sigma^2}}$ is the exponential of a kernel function, and as a consequence a kernel function itself

5. $e^{-\frac{\mathbf{x}_1^T \mathbf{x}_1}{\sigma^2}} e^{-\frac{\mathbf{x}_2^T \mathbf{x}_2}{2\sigma^2}} e^{\frac{\mathbf{x}_1^T \mathbf{x}_2}{\sigma^2}}$ is a kernel function, being the product of a

kernel function with two functions $f(\mathbf{x}_1) = e^{-\frac{\mathbf{x}_1^T \mathbf{x}_1}{\sigma^2}}$ and

Relevant kernel functions

- Polynomial kernel

$$\kappa(\mathbf{x}_1, \mathbf{x}_2) = (\mathbf{x}_1 \cdot \mathbf{x}_2 + 1)^d$$

- Sigmoidal kernel

$$\kappa(\mathbf{x}_1, \mathbf{x}_2) = \tanh(c_1 \mathbf{x}_1 \cdot \mathbf{x}_2 + c_2)$$

- Gaussian kernel

$$\kappa(\mathbf{x}_1, \mathbf{x}_2) = \exp\left(-\frac{\|\mathbf{x}_1 - \mathbf{x}_2\|^2}{2\sigma^2}\right)$$

where $\sigma \in \mathbb{R}$

Observe that a gaussian kernel can be derived also starting from a non linear kernel function $\kappa(\mathbf{x}_1, \mathbf{x}_2)$ instead of $\mathbf{x}_1^T \mathbf{x}_2$.

Kernels of structured objects

Kernels are particularly useful when applied to structured objects.

Consider the case of strings (for example sequences of DNA bases or amino acids).

Given two strings $\mathbf{x}_1, \mathbf{x}_2$ on a same alphabet \mathcal{A} , we can define their similarity to be equal to the number of substrings they have in common.

More formally, let $\phi_s(\mathbf{x})$ be the number of times substring s occurs in \mathbf{x} and let $\boldsymbol{\phi}(\mathbf{x})$ the corresponding vector of such functions for all substrings s : a kernel can be defined as

$$\kappa(\mathbf{x}_1, \mathbf{x}_2) = \boldsymbol{\phi}(\mathbf{x}_1)^T \boldsymbol{\phi}(\mathbf{x}_2) = \sum_{s \in \mathcal{A}^*} w_s \phi_s(\mathbf{x}_1) \phi_s(\mathbf{x}_2)$$

where $w_s \geq 0$ are predefined weights.

If $w_s = 1$ for all considered substrings and we define $\phi'(\mathbf{x}) = \frac{\phi(\mathbf{x})}{\|\phi(\mathbf{x})\|^2}$ as a normalized version of ϕ , we get

$$\kappa(\mathbf{x}_1, \mathbf{x}_2) = \frac{\phi(\mathbf{x}_1)^T \phi(\mathbf{x}_2)}{\|\phi(\mathbf{x}_1)\|^2 \|\phi(\mathbf{x}_2)\|^2}$$

that is, the well known cosine similarity measure.

Borrowing from information retrieval methods, a better similarity measure can be obtained by defining $\phi_s(\mathbf{x})$ through more sophisticated measures, such as tf-idf, instead of occurrences counting

Special cases:

- $w_s = 0$ if $|s| > 1$ is a bag-of-chars kernel, with $\phi_c(\mathbf{x})$ being the number of occurrences of character c in \mathbf{x}
- If only s delimited by white spaces are considered, we get a bag-of-words kernel.
- If only strings of fixed length $|s| = k$ are considered, we have a k -spectrum kernel
- ...

The approach can be extended to the case of trees, in order to deal with, for example, parse or evolutionary trees

More complex kernel construction techniques have been defined.